Build plan — 4 sessions, checkpoint + image evidence after each

Rule: every session ends with (a) a working, runnable checkpoint and (b) a screenshot saved under `docs/img/`. No session is "done" without the image.

#	Session	Builds	Image evidence
1	**Real terrain field**	Pull a real DEM of Lake San Antonio → derive sector graph (water/steep = gated ravine, edge cost = real slope) → **swap the abstract grid in the working sandbox for the real field**. Add the resource readout + 90-min phase clock.	live match running on the real hillshade, with sectors/bridges/squads
2	**Game-design mechanics**	Implement the design doc: **classes** (capture weights/abilities), **intel/fog** (see only scouted sectors), **resource sinks** (bridge/intel/siege), and the **phase-gated 90-min arc** (land-grab → contact → siege bell).	fogged map + class roster + phase banner + resource ledger
3	**RL playtesters**	Wrap the engine as a gym/PettingZoo env with a **microrts-style observation tensor + command action**; train **PPO self-play** at small scale (3v3, ~14 sectors); export trained-agent episodes.	learning curve climbing + an RL episode replay + the strategy it found
4	**Balance verdict + polish**	**OpenSpiel-style** metrics (fairness, matchup matrix, exploitability), wire RL + heuristics into the Monte Carlo, link findings to the honesty ledger, polish UI, finalize the design doc.	the full "money-shot" screen with **real data** (not a mockup)

Each checkpoint is committable and shippable on its own — if we stop after any session, there's still a coherent artifact.

Known risk (Session 1): fetching real elevation data depends on network/tool access in the build environment. Primary source: USGS 3DEP via OpenTopography; fallback: AWS Terrain Tiles (Terrarium PNG → decode). If both are blocked, I'll use a real DEM file you drop in `data/` — the pipeline is identical either way.